

Working with Oracle SQL

Chapter 6:

Data Definition Language to create and manage tables

Chapter Objectives

In this chapter, we will discuss:

- Creating and managing tables
- Altering tables
- Using sequence generators
- Managing integrity constraints
- Creating and managing views
- Working with indexes



Schemas

Create and Manage Tables

Alter a Table

Creating Sequence Generators

Enforcing Database Integrity

Managing Views

Managing Indexes

Using SQL Developer

Chapter Summary

Schema

- “Schema” is a relational database word meaning:
 - Collection of database objects
 - Including structures such as:
 - Tables
 - Views
 - Stored procedures
 - Indexes
- A schema has the name of the user who owns it
 - The terms schema, user, and account tend to be used interchangeably
 - HR and SCOTT are the users/schemas we have been using

Schema Objects

- Everything in the database *must* be owned by a user
 - The collection of objects owned by a user is called the schema
- A schema object is owned by the user account it is created in
 - Qualified by: `user_name.table_name`
 - Example: `hr.employees`
- Legal schema object names must conform to the Oracle standard
 - Maximum of 30 positions
 - Must start with an alpha character
 - Can include numbers and the special characters: `$ _ #`

Schemas



Create and Manage Tables

Alter a Table

Creating Sequence Generators

Enforcing Database Integrity

Managing Views

Managing Indexes

Using SQL Developer

Chapter Summary

CREATE TABLE

- The simplified syntax is:

```
CREATE TABLE table_name (  
    column_name datatype [DEFAULT default_value] [NOT NULL]  
    [, ... ]  
)
```

- A table must have at least one column
 - Each column must have a valid datatype
 - Can be defined as being mandatory
 - Can have a `DEFAULT` value set
- `CREATE TABLE` is a **DDL** command
 - DDL commands automatically commit
 - Any DDL command automatically commits any outstanding transaction

CREATE TABLE Example

- A table to store client information

```
CREATE TABLE clients (  
    client_id          NUMBER(5) NOT NULL  
    , client_name      VARCHAR2(20)  
    , client_type      NUMBER(2) DEFAULT 01  
    , client_start_date DATE  
);  
  
Table created.
```

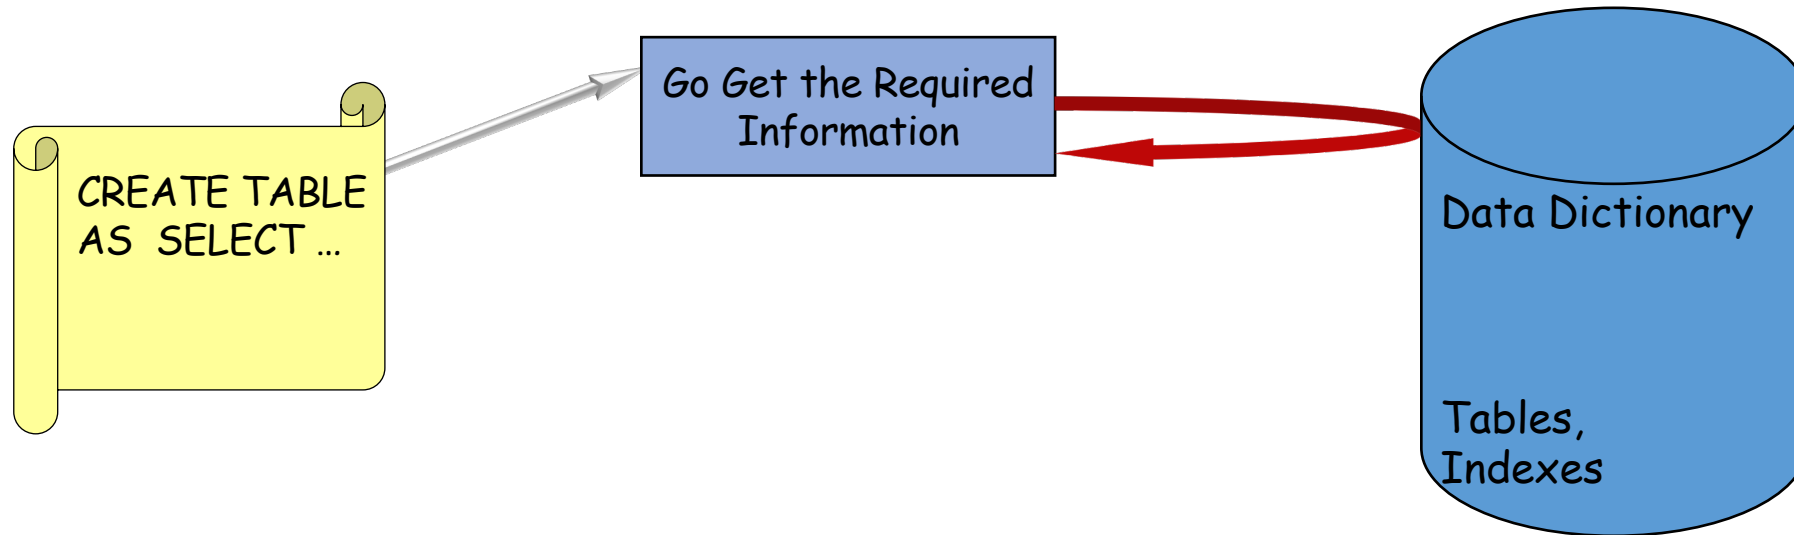
- NOT NULL is an integrity constraint specifying that the column is mandatory
- The client type is set to 01, if a value is not provided when a row is inserted
 - Using the key word DEFAULT in an UPDATE or INSERT will cause this value to be used

CREATE TABLE AS SELECT: CTAS

- IF you want to base the definition of a new table upon an existing one, the CTAS statement can be used
- Syntax:

```
CREATE TABLE table_name AS SELECT ....
```

- The necessary column specifications are retrieved from the data dictionary



CTAS: Example 1

- An image copy of an existing table can be made
 - Notice that the rows of data are brought over

```
CREATE TABLE new_jobs AS
  SELECT *
  FROM jobs;
```

Table created.

```
SELECT *
FROM new_jobs;
```

JOB_ID	JOB_TITLE	MIN_SALARY	MAX_SALARY
AD_PRES	President	20000	40000
AD_VP	Administration Vice President	15000	30000
AD_ASST	Administration Assistant	3000	6000
FI_MGR	Finance Manager	8200	16000
...			

19 rows selected.

CTAS: Example 1 (Continued)

- IF you wanted a copy of the definition of the table without any data, write the `WHERE` clause to exclude all data

```
CREATE TABLE new_jobs AS  
  SELECT *  
  FROM jobs  
  WHERE 1 = 2;
```

Table created.

```
SELECT * FROM new_jobs;
```

No rows selected.

CTAS: Example 2

- We can also specify which columns we want and their names by providing a column list

```
CREATE TABLE new_jobs (title, job_number) AS
  SELECT job_title, job_id
  FROM jobs;
```

Table created.

```
SELECT * FROM new_jobs;
```

TITLE	JOB_NUMBER
-----	-----
President	AD_PRES
Administration Vice President	AD_VP
Administration Assistant	AD_ASST
Public Relations Representative	PR_REP
. . .	

19 rows selected.

CTAS: Example 3

- We can also limit the rows in the new table by specifying a `WHERE` clause

```
CREATE TABLE new_jobs (title, job_number) AS
  SELECT job_title, job_id
  FROM jobs
  WHERE job_id LIKE 'IT%';
```

Table created.

```
SELECT * FROM new_jobs;
```

TITLE	JOB_NUMBER
Programmer	IT_PROG

CTAS: Example 4

- The table being created does not have to be based upon a single table
 - Any `SELECT` statement can be used

```
CREATE TABLE jobs_employees AS
  SELECT j.job_title, e.last_name
  FROM jobs j
  JOIN employees e
  USING (job_id)
  WHERE job_id LIKE 'IT%';
```

Table created.

```
SELECT * FROM jobs_employees;
```

JOB_TITLE	LAST_NAME
Programmer	Hunold
Programmer	Ernst
Programmer	Austin
Programmer	Pataballa
Programmer	Lorentz

Other Table Commands

- Tables can be renamed

```
RENAME table_name TO new_table_name  
RENAME jobs TO regional_jobs;
```

- Be careful:
 - Permissions and references to the original table are not automatically maintained
- Tables can have their rows quickly removed

```
TRUNCATE TABLE new_jobs;
```

 - Accomplishes the same result as `DELETE` without the `WHERE` clause
- Be careful:
 - `TRUNCATE` is a DDL command and cannot be rolled back
 - The reason why it is so efficient

Dropping Tables

- To remove the definition of the table (not just the rows):

```
DROP TABLE table_name  
DROP TABLE new_jobs;
```

- This removes the table and all dependencies
 - Indexes, permissions, etc.
 - Code references become invalid

Restoring Dropped Tables

- Beginning in Oracle10g, the table is not actually dropped
 - The table is renamed and “held” in the recycle bin
 - Along with dependent objects, indexes, triggers, constraints
- The table can be restored using a FLASHBACK command
 - `FLASHBACK TABLE table_name TO BEFORE DROP`
 - `FLASHBACK new_jobs TO BEFORE DROP;`
 - Dependent objects, except for foreign keys, are also restored
- The table can also be renamed as a part of the flashback operation
 - `FLASHBACK new_jobs TO BEFORE DROP RENAME TO some_jobs;`
- Unless purged, the contents of the recycle bin will be available as long as space exists in the area where the original table was stored
- The recycle bin can be bypassed during the DROP by appending the PURGE option
 - `DROP TABLE new_jobs PURGE;`

Working with the Recycle Bin

- To examine the recycle bin

```
SELECT ORIGINAL_NAME, OBJECT_NAME, TYPE, DROPTIME, CAN_UNDROP FROM user_recyclebin;
```

ORIGINAL_NAME	OBJECT_NAME	TYPE	DROPTIME	CAN_UNDROP
NEW_JOBS	BIN\$Y/ym7LmtSh2IYYBGAoBqoQ==\$0	TABLE	2018-05-03:09:36:53	YES

- Can access objects directly from the bin

```
SELECT * FROM "BIN$Y/ym7LmtSh2IYYBGAoBqoQ==$0";
```

- A specific table can be removed from the recycle bin

- Or, if there was more than one matching table:

```
PURGE TABLE new_jobs;
```

```
PURGE TABLE "BIN$Y/ym7LmtSh2IYYBGAoBqoQ==$0";
```

- The entire contents of the recycle bin can be purged with a single command

- Verify the contents first
- Requires DBA privileges

```
PURGE DBA_RECYCLEBIN;
```

Chapter Concepts

Schemas

Create and Manage Tables



Alter a Table

Creating Sequence Generators

Enforcing Database Integrity

Managing Views

Managing Indexes

Using SQL Developer

Chapter Summary

Altering Tables

- A table, once created, can have many of its settings changed via the `ALTER` command
- For example, to add a new column to an existing table:

```
ALTER TABLE table_name
ADD (column_name datatype [DEFAULT default_value] [NOT NULL]
    [, ...] )
```

```
ALTER TABLE new_jobs
ADD (effective_date DATE);
Table altered.
```

```
DESCRIBE new_jobs
```

Name	Null?	Type
-----	-----	-----
JOB_ID		VARCHAR2(10)
JOB_TITLE	NOT NULL	VARCHAR2(35)
MIN_SALARY		NUMBER(6)
MAX_SALARY		NUMBER(6)
EFFECTIVE_DATE		DATE

Can only add a NOT NULL column if the table is empty, or there is a default specified

Altering Tables: Modifying Columns

- Using a similar syntax, existing columns can be modified
- Example:

```
ALTER TABLE new_jobs  
MODIFY (max_salary NUMBER(8));
```

Table altered.

```
DESCRIBE new_jobs
```

Name	Null?	Type
-----	-----	-----
JOB_ID		VARCHAR2 (10)
JOB_TITLE	NOT NULL	VARCHAR2 (35)
MIN_SALARY		NUMBER (6)
MAX_SALARY		NUMBER (8)
EFFECTIVE_DATE		DATE

Restrictions on Modifying Columns

- If rows already exist, then the pre-existing data must already agree with the modification
- Examples:
 - A column's size cannot be made smaller than the existing data
 - A column's datatype cannot be changed as long as there is existing data in the column
 - Any existing constraints must be adhered to

```
ALTER TABLE new_jobs  
MODIFY (job_title VARCHAR2(4));  
(job_title VARCHAR2(4) )  
*
```

ERROR at line 2:

ORA-01441: cannot decrease column length because some value is too big

```
ALTER TABLE new_jobs  
MODIFY (job_id NUMBER(6));  
(job_id NUMBER(6) )  
*
```

ERROR at line 2:

ORA-01439: column to be modified must be empty to change datatype

Altering Tables: Removing Columns

- Columns can be removed from the table in one of several ways
 - Remove the definition and the data

```
ALTER TABLE new_jobs
DROP (min_salary, max_salary) CASCADE CONSTRAINTS;
```

Table altered.

```
DESCRIBE new_jobs
```

Name	Null?	Type
-----	-----	-----
JOB_ID		VARCHAR2(10)
JOB_TITLE	NOT NULL	VARCHAR2(35)
EFFECTIVE_DATE		DATE

- The **CASCADE** clause is needed when the column is a Foreign Key
 - In this example, the clause is unnecessary

Altering Tables: Setting Columns Unused

- The data dictionary can be modified without physically removing the column's data
 - The SET UNUSED clause

```
ALTER TABLE new_jobs  
SET UNUSED (effective_date) CASCADE CONSTRAINTS;
```

Table altered.

```
DESCRIBE new_jobs
```

Name	Null?	Type
-----	-----	-----
JOB_ID		VARCHAR2(10)
JOB_TITLE	NOT NULL	VARCHAR2(35)

- Allows the data to be physically removed when convenient

```
ALTER TABLE new_jobs  
DROP UNUSED COLUMNS;  
Table altered.
```


Altering Tables: Renaming Columns

- Existing columns can be renamed with the `RENAME` clause

```
ALTER TABLE new_jobs  
RENAME COLUMN job_title TO title;
```

Table altered.

```
DESCRIBE new_jobs
```

Name	Null?	Type
JOB_ID		VARCHAR2 (10)
TITLE	NOT NULL	VARCHAR2 (35)

- Any dependent constraints are maintained
- Dependent code (triggers, views, stored procedures, etc.) are invalidated

Chapter Concepts

Schemas

Create and Manage Tables

Alter a Table



Creating Sequence Generators

Enforcing Database Integrity

Managing Views

Managing Indexes

Using SQL Developer

Chapter Summary

Sequence Generator

- A sequence generator is a database object
 - It is an Oracle-managed series of numbers
 - Often used to provide the values for primary keys
 - Remember, primary keys must be unique
- Shortened (simplified) syntax:

```
CREATE SEQUENCE sequence_name  
    INCREMENT BY interval  
    START WITH start_value
```

- Example:

```
CREATE SEQUENCE seq_emp  
    INCREMENT BY 1  
    START WITH 8000;
```

Sequence created.

Sequence Generator Example

- We can now use this sequence to provide the data value for the primary key (`empno`) of the `emp` table
- Get the next sequence value by using `NEXTVAL`

```
INSERT INTO emp (empno, ename)
VALUES (seq_emp.NEXTVAL, 'CHAVAS');

1 row created.
```

- Now retrieve the rows

```
SELECT empno, ename FROM emp WHERE empno > 7900;
   EMPNO  ENAME
-----
    7902  FORD
    7934  MILLER
    8000  CHAVAS
```

- A sequence number, once consumed, cannot be put back (even if rollback occurs)

IDENTITY Columns

- Sequences are normally combined with triggers to populate key fields automatically
 - Starting with Oracle 12c, IDENTITY columns make this functionality more accessible

```
GENERATED [ ALWAYS | BY DEFAULT [ ON NULL ] ] AS IDENTITY [ ( options ) ]
```

- ALWAYS means it will always be auto-generated, specifying a value is an error
- BY DEFAULT means it will be auto-generated unless specified
- BY DEFAULT ON NULL means auto-generated if set to NULL
- options are the same as for a SEQUENCE

```
CREATE TABLE clients (  
    client_id          NUMBER GENERATED ALWAYS AS IDENTITY  
    , client_name      VARCHAR2(20)  
    , client_type      NUMBER(2) DEFAULT 01  
    , client_start_date DATE  
);
```

Table created.

Chapter Concepts

Schemas

Create and Manage Tables

Alter a Table

Creating Sequence Generators



Enforcing Database Integrity

Managing Views

Managing Indexes

Using SQL Developer

Chapter Summary

Integrity Constraints

- Declarative method of enforcing business rules in the database
- Defined with the corresponding table using the `CREATE TABLE` or `ALTER TABLE` command
- Validated when constraints are defined and on each `INSERT`, `UPDATE`, and `DELETE`
- Five types
 - `NOT NULL`
 - Primary key
 - Unique key
 - Foreign key
 - Check

NOT NULL Constraints

- Enforces that a column contains a value for each row
- Usually defined in the `CREATE TABLE` command, as discussed in the previous section
- Example:
 - `empno` is mandatory

```
CREATE TABLE emp (  
    empno NUMBER(4)      NOT NULL  
    , ename VARCHAR2(10)  
    ,  ...  
);
```

- May be used with `ALTER TABLE MODIFY` after adding a new column to a populated table
 - Add column without the constraint
 - Populate the column
 - Alter the table to add the constraint

Primary Key Constraints

- Ensures that no two rows of a table have duplicate values in the key column(s)
 - All columns in the primary key must be mandatory
 - Only one primary key is allowed per table
- For ease of maintenance, use the `ALTER TABLE` command with the following syntax:

```
ALTER TABLE table_name  
ADD CONSTRAINT constraint_name  
PRIMARY KEY (column_name, ..., column_name);
```

Primary Key Constraints (continued)

- An ideal primary key column is one whose value rarely changes
 - Updating a primary key usually impacts other tables
- A common naming convention is to end primary keys with the `_pk` suffix
 - Makes it easier to identify these keys when we query the data dictionary
- Example:
 - `empno` is used to uniquely identify employees

```
ALTER TABLE emp  
ADD CONSTRAINT emp_pk  
PRIMARY KEY (empno);
```

- Short numeric columns are preferred as primary keys because they offer the best performance and require the least amount of storage

Unique Key Constraints

- Same rules as for primary key, but allow optional columns
 - Allow an unlimited number per table
 - A common standard is to end unique keys with the `_uk` suffix
- For ease of maintenance, use the `ALTER TABLE` command with the following syntax:

```
ALTER TABLE table_name  
ADD CONSTRAINT constraint_name  
UNIQUE (column_name, ..., column_name);
```

- Example:
 - `ename, job, and hiredate` must be unique

```
ALTER TABLE emp  
ADD CONSTRAINT emp_uk1  
UNIQUE (ename, job, hiredate);
```

Foreign Key Constraints

- Ensures that every value of a foreign key in the child table exists in the parent table
 - Prevents updates and deletes of parent table that would violate the constraint

```
ALTER TABLE child_table_name
ADD CONSTRAINT constraint_name
FOREIGN KEY (column_name, ..., column_name)
REFERENCES parent_table_name [(column_name, ..., column_name)]
[ ON DELETE CASCADE | ON DELETE SET NULL ];
```

- A foreign key can only reference existing primary or unique keys of the parent table
 - If no parent columns are specified, references primary key
- **The ON DELETE options:**
 - CASCADE automatically deletes child records on deletion of parent record
 - SET NULL automatically updates child columns to NULL on deletion of parent record
- A common naming convention is to end foreign keys with the `_fk` suffix

Foreign Key Constraint Examples

- Each employee must belong to a department that is defined in the `dept` table
 - A department cannot be deleted if employees work there

```
ALTER TABLE emp
ADD CONSTRAINT emp_dept_fk1
FOREIGN KEY (deptno) REFERENCES dept (deptno);
```

- Each employee must belong to a department that is defined in the `dept` table
 - If a department is being deleted, all of its employees should be deleted

```
ALTER TABLE emp
ADD CONSTRAINT emp_dept_fk1
FOREIGN KEY (deptno) REFERENCES dept (deptno)
ON DELETE CASCADE;
```



Make sure this is
really what you
want!

Check Constraints

- Enforces business rules on update or insert
 - Limited to simple validation based on constants or column values in the row

```
ALTER TABLE table_name  
ADD CONSTRAINT constraint_name  
CHECK (condition);
```

- A common naming convention is to end check constraints with the `_ck` suffix
- E.g., Commission must always be less than 10% of salary

```
ALTER TABLE emp  
ADD CONSTRAINT emp_comm_sal_ck  
CHECK (comm < sal * .10);
```

Dropping Constraints

- Drop constraints using the following syntax:

```
ALTER TABLE table_name  
DROP CONSTRAINT constraint_name;
```

- For example, drop emp_comm_sal_ck check constraint

```
ALTER TABLE emp  
DROP CONSTRAINT emp_comm_sal_ck;
```

- A simpler way to drop primary key constraints is using this syntax:

```
ALTER TABLE table_name  
DROP PRIMARY KEY [CASCADE];
```

- CASCADE option drops all the related foreign key constraints
- For example, drop the primary key on emp table

```
ALTER TABLE emp  
DROP PRIMARY KEY;
```

Enabling and Disabling Constraints

- Constraints can be made temporarily inactive and can later be reactivated using the following syntax:

```
ALTER TABLE table_name  
DISABLE CONSTRAINT constraint_name;
```

```
ALTER TABLE table_name  
ENABLE CONSTRAINT constraint_name;
```

- Examples:

- Disable emp_dept_fk1 foreign key constraint, then re-enable it

```
ALTER TABLE emp DISABLE CONSTRAINT emp_dept_fk1;  
  
Table EMP altered  
  
ALTER TABLE emp ENABLE CONSTRAINT emp_dept_fk1;  
  
Table EMP altered
```


Constraints in the Dictionary

- Constraints are shown in the `user_constraints` dictionary view
 - Constraint type is “P” for primary key, “U” for unique key, “R” for referential constraint or foreign key, “C” for check constraint, and “V” for view with check option
 - Constraint status is either `ENABLED` or `DISABLED`
 - Constraint columns are stored in `user_cons_columns` view
- Example:

```
SELECT constraint_name, constraint_type, status
FROM user_constraints
WHERE table_name = 'EMP';
```

CONSTRAINT_NAME	C	STATUS
-----	-	-----
EMP_PK	P	ENABLED
EMP_DEPTNO_FK1	R	ENABLED

Chapter Concepts

Schemas

Create and Manage Tables

Alter a Table

Creating Sequence Generators

Enforcing Database Integrity



Managing Views

Managing Indexes

Using SQL Developer

Chapter Summary

Views

- A view is a stored definition of data
 - The data dictionary is updated when a view is created
 - The data is not physically stored in a view
 - It is read from the tables the view references

- Syntax:

```
CREATE OR REPLACE VIEW view_name (column_name, column_name ...) AS  
SELECT ...  
[WITH { READ ONLY } | { CHECK OPTION [CONSTRAINT constraint_name ] } ]
```

- The `SELECT` statement can be any valid query
- The view can specify the column names
- There are two possible restrictions
 - `READ ONLY` means the view data cannot be updated
 - `CHECK OPTION` restricts updates to values that would be selected by the view

Reasons to Create Views

- Security
 - To restrict the columns that can be seen
 - To restrict the rows that can be read
- To simplify the accessing of information
 - Complex SQL statements do not have to be continually recreated
- To provide typical application images of the information
 - Including using a different column name
 - Preferred by different groups
- Isolate applications from table changes

VIEW Example

- Provide a restricted view of the jobs table for other accounts
 - Only three of the columns and some of the rows in the jobs table can be seen
 - So that a restricted user can only access the data in the view

```
CREATE OR REPLACE VIEW vw_jobs (max_sal, min_sal, title) AS
SELECT max_salary, min_salary, job_title
FROM jobs
WHERE max_salary < 8000
WITH CHECK OPTION CONSTRAINT ck_restrict_jobs;
```

View created.

```
SELECT * FROM vw_jobs;
   MAX_SAL   MIN_SAL  TITLE
-----
      6000       3000 Administration Assistant
      5500       2500 Purchasing Clerk
      5000       2000 Stock Clerk
      5500       2500 Shipping Clerk
```

Updatable Views

- Views may be updatable if they contain a key-preserved table
 - A table where each row appears only once in the view results
- Each operation on an updatable view may only modify the data of one underlying table

```
UPDATE vw_jobs  
SET max_sal = 7500  
WHERE title = 'Administration Assistant';
```

1 row updated.

```
UPDATE vw_jobs  
SET max_sal = 8500  
WHERE title = 'Administration Assistant';
```

Error report -

```
ORA-01402: view WITH CHECK OPTION where-clause violation
```

Updatable Views (continued)

- To find out if a view column is updatable

```
SELECT table_name, column_name, updatable
FROM   user_updatable_columns
WHERE  table_name = 'VW_JOBS';
```

TABLE_NAME	COLUMN_NAME	UPD
-----	-----	---
VW_JOBS	MAX_SAL	YES
VW_JOBS	MIN_SAL	YES
VW_JOBS	TITLE	YES

Chapter Concepts

Schemas

Create and Manage Tables

Alter a Table

Creating Sequence Generators

Enforcing Database Integrity

Managing Views



Managing Indexes

Using SQL Developer

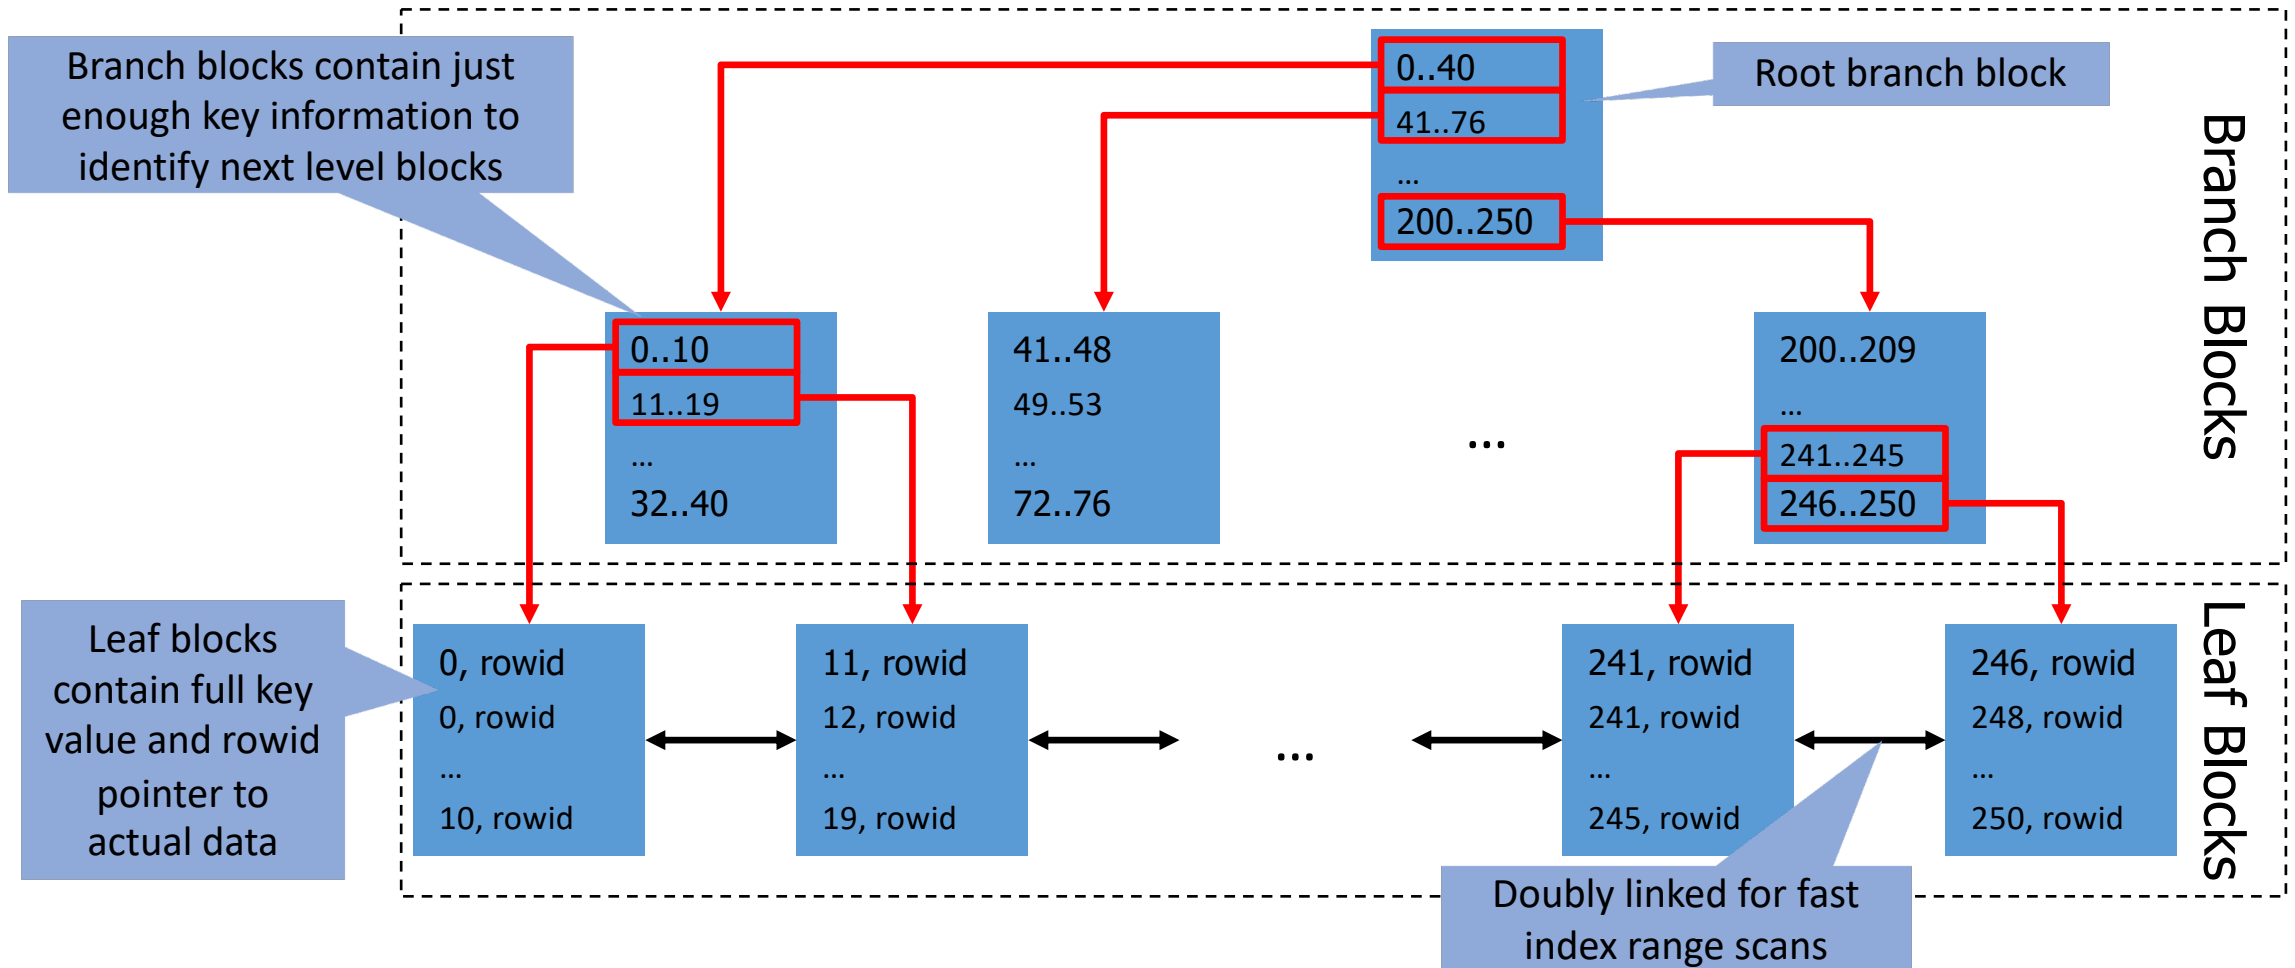
Chapter Summary

Indexes

- An index is a schema object for improving query performance
 - Stored separately from the table
 - Not referenced in SQL, used by query optimizer
- Indexes work because keys are generally much smaller than the data they reference
 - Index stores only the key and a pointer to the actual data (rowid)
 - Index can be traversed much faster than the data
 - Oracle automatically creates indexes for primary key and unique constraints
- Default Oracle index type is B-Tree index
 - “Balanced” tree because each leaf node is the same number of steps away from the root
 - Accessing leaf data requires $O(\log N)$ block accesses (increases very slowly)

B-Tree Indexes

- Imagine a numeric index, e.g., foreign key index for `department_id` in `employees`



Unique Index

- An index with a one-to-one relationship to table rows
 - Efficient because each leaf entry points to one row in the table
 - Automatically generated for each primary key and unique key constraint
 - Unless there is already a covering index
 - They are dropped when the underlying constraint is dropped or disabled

```
ALTER TABLE emp  
ADD CONSTRAINT emp_pk  
PRIMARY KEY (empno);
```

```
ALTER TABLE emp  
ADD CONSTRAINT emp_uk1  
UNIQUE (ename, hiredate);
```

- Or you can create your own

```
CREATE UNIQUE INDEX emp_pk_idx  
ON emp (empno);
```

```
CREATE UNIQUE INDEX emp_uk1_idx  
ON emp (ename, hiredate);
```

- Indexes are frequently named with a suffix `_ix`

Non-Unique Index

- A non-unique index has values that point to one or more table rows
 - More than one rowid per key value
 - They are generally less efficient than unique indexes
 - How much less efficient depends on the selectivity of the columns

```
CREATE INDEX index_name ON table_name (column_name, ..., column_name)
```

- Example:
 - Manager in the emp table is not unique
 - More than one employee can work for the same manager
 - If we frequently asked the question “who works for...”, we might create an index to improve performance

```
CREATE INDEX emp_mgr_ix  
ON emp (mgr);
```

Dropping Indexes

- Any index may be dropped unless it is being used to enforce a constraint

```
DROP INDEX emp_mgr_ix;
```

- Automatically created indexes for primary keys and unique constraints are automatically dropped if the constraint is dropped
 - You cannot usually drop them manually
- Any index you create may be used by a constraint that is created later
 - In this case, you will not be able to drop the index without first dropping the constraint

Function-Based Indexes

- You can create an index that includes a function call
 - SQL scalar function
 - User-defined function
 - Package function
- Oracle pre-computes the function return value and stores it in the index

```
CREATE INDEX emp_fname_uppercase_ix  
ON employees (UPPER(first_name));
```

- What do you think the performance impact of this will be:
 - For queries?
 - For data maintenance (insert and update)?

Indexing Guidelines

- Deciding WHAT indexes to create is driven by HOW the data is accessed
 - This can change over time
- Columns to consider for an index
 - All columns that are part of primary and unique keys are indexed automatically
 - Foreign keys that are commonly used for joins
 - Frequently queried columns that are relatively unique (high selectivity)
 - Return less than 10 percent of the rows when queried
- Columns not to index
 - Columns that are frequently updated
 - Columns in small tables
 - Columns that contain few unique values (low selectivity)
 - Reason: likely to be retrieving many rows

Exercise 6.1: Table Management



45 min

- Please turn to the Exercise Manual and complete this exercise

Chapter Concepts

Schemas

Create and Manage Tables

Alter a Table

Creating Sequence Generators

Enforcing Database Integrity

Managing Views

Managing Indexes

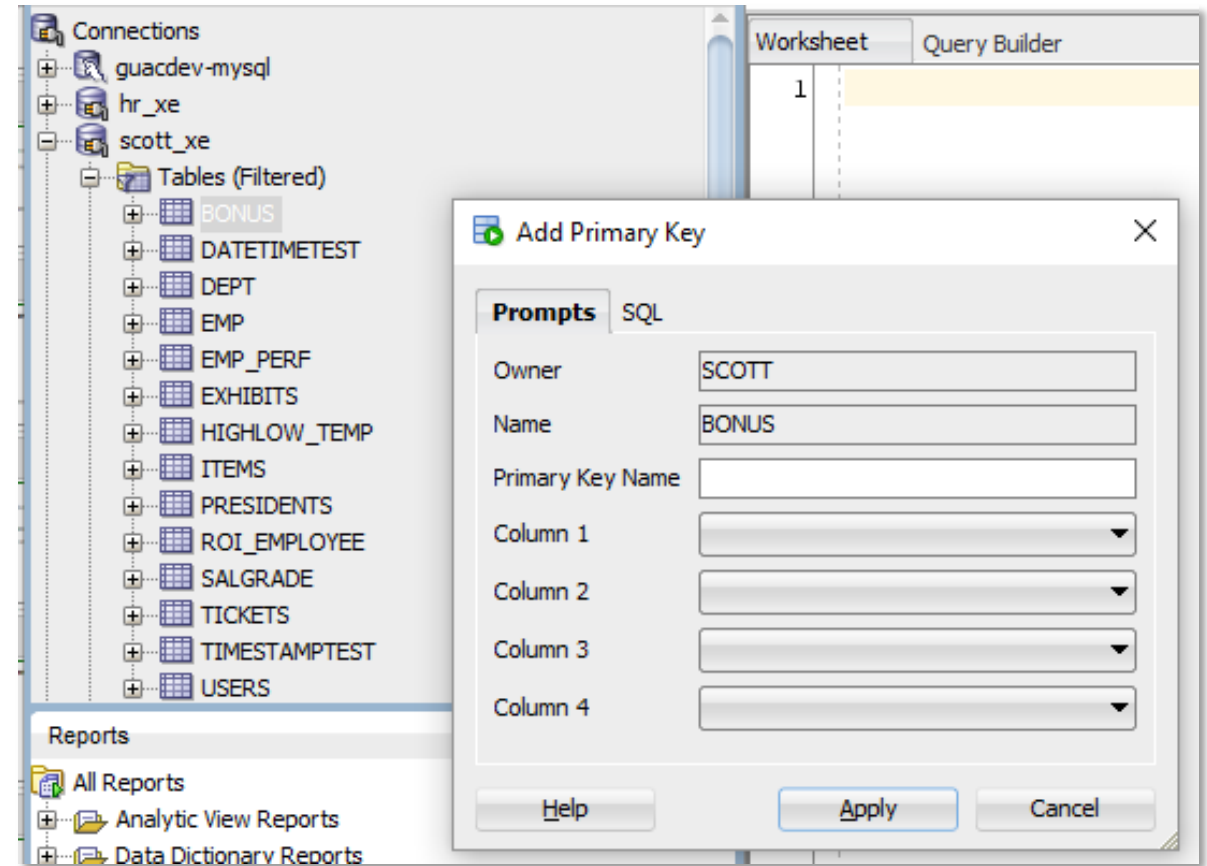


Using SQL Developer

Chapter Summary

DDL Through SQL Developer

- SQL Developer provides a method of doing most DDL through the GUI
- ***Do not use this feature***
- Get in the habit of typing the commands:
 - It helps you learn
 - It means there is a clear history
 - It allows you to maintain a DDL file
- If you are unsure of the syntax, use SQL Developer and then immediately export the DDL to examine it and to keep a history



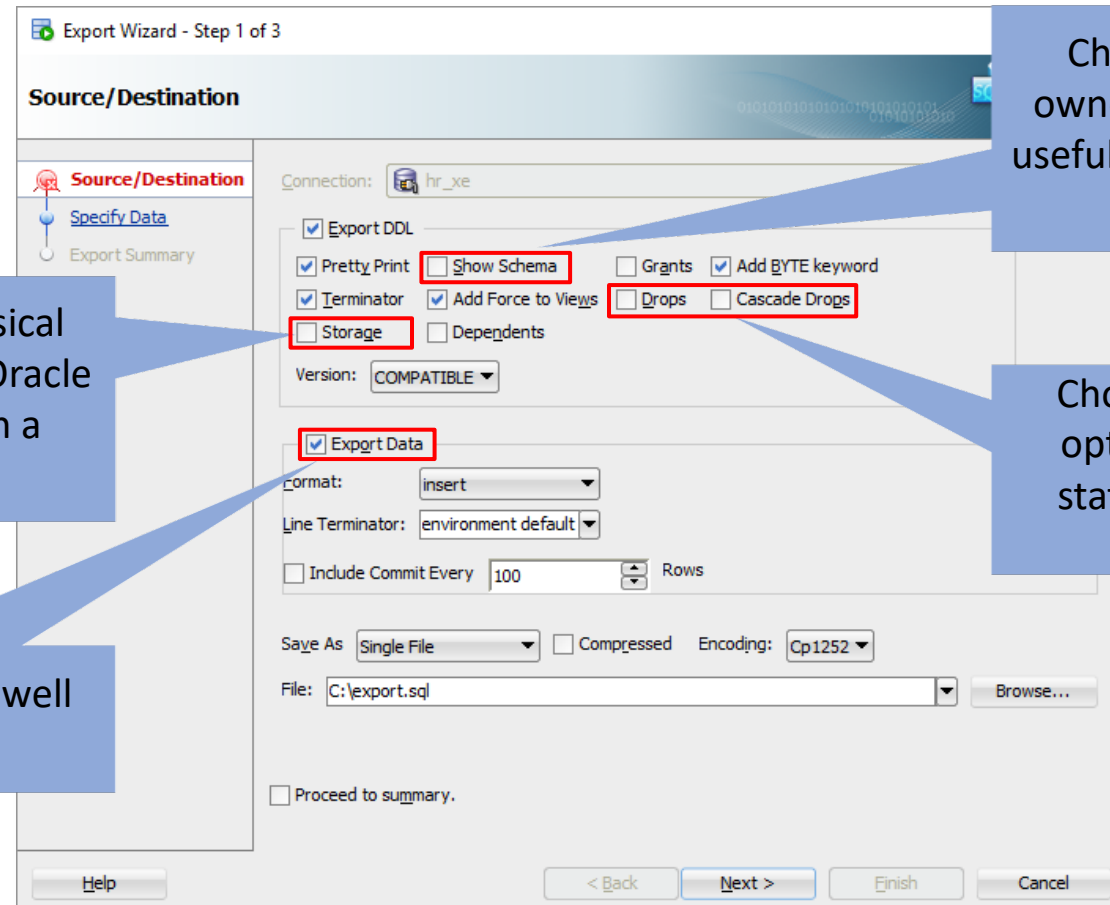
Maintaining DDL Files

- You should maintain a series of DDL files for your project
 - You should be able to completely recreate the database from these scripts
 - Separate test data from schema structure
 - Consider separating schema objects into separate files and including them using @@

```
@@client_ddl.sql  
@@staff_ddl.sql  
  
@@populate_data.sql
```

Exporting from SQL Developer

- You can generate a starting point by exporting from SQL Developer
 - Experiment with the options



Checking this option shows schema ownership for each object, which is not useful if you want to re-create the objects in another schema

Adds information about the physical storage that you usually leave to Oracle and makes it harder to apply on a different server

Choosing one of these options creates DROP statements before the CREATEs

If you want the data as well as the structure

Chapter Concepts

Schemas

Create and Manage Tables

Alter a Table

Creating Sequence Generators

Enforcing Database Integrity

Managing Views

Managing Indexes

Using SQL Developer



Chapter Summary

Chapter Summary

In this chapter, we have discussed:

- Creating and managing tables
- Altering tables
- Using sequence generators
- Managing integrity constraints
- Creating and managing views
- Working with indexes