

Working with Oracle SQL

Chapter 7: Data Manipulation Language

Chapter Objectives

In this chapter, we will discuss:

- Manipulating data using DML:
 - INSERT
 - UPDATE
 - DELETE
- Transactional Control Statements:
 - COMMIT
 - ROLLBACK
 - SAVEPOINT



INSERT

UPDATE

DELETE

Transactional Control

Chapter Summary

Data Manipulation Language Statements

■ Data Manipulation Language (DML) is a classification of SQL

– Others we are talking about:

- Data Query Language (DQL)
 - The `SELECT` statement
- Data Definition Language (DDL)
 - `CREATE`, `ALTER`
- Transaction Control Language (TCL)
 - `COMMIT`, `ROLLBACK`, `SAVEPOINT`

■ DML statements control the values of the data

Adding Data: the INSERT Statement

- DML statements allow you to manipulate the data in any table in your account
 - Or to have permission, if the table is owned by another user
 - If synonyms are not defined, qualify the table name: `schema.table`
- INSERT, UPDATE, and DELETE statements operate against a set of data
 - A set is zero, 1, or many rows
- The syntax for the INSERT allows either the insertion of one row or an entire set of rows

INSERTing One Row of Data

■ Syntax:

```
INSERT INTO table_name [ (column_list) ]  
VALUES (value_clause)
```

■ Rules:

- The values list must map to the column list one for one
- All integrity rules must be met
- If there is a defined default value for a column, it is invoked using the key word: `DEFAULT`
 - This table setting is covered later
- The column list is optional
 - Values list must include a value for each column in the table in the order they appear in the table
 - Even if the value is `NULL`
 - Only use this for ad hoc inserts; always use the column list in production code

INSERTing One Row of Data: Examples

■ The following are equivalent:

```
INSERT INTO regions  
VALUES (5, 'Antarctica');
```

```
INSERT INTO regions (region_id, region_name)  
VALUES (5, 'Antarctica');
```

```
INSERT INTO regions (region_name, region_id)  
VALUES ('Antarctica', 5);
```

```
INSERT INTO regions  
VALUES (5/1 ,SUBSTR('Antarctica', 1));
```

INSERTing One Row of Data: More Examples

■ Unless a column is mandatory, a value does not have to be inserted

- Columns for which no value is specified are set to `NULL`
- Or to their `DEFAULT` value

```
INSERT INTO regions (region_id) VALUES (5);
```

1 row created.

```
SELECT *  
FROM regions  
WHERE region_id = 5;
```

REGION_ID	REGION_NAME
5	

Accepting Default Values for Columns

■ IF the `region_name` column had a `DEFAULT` value, then it would be used if a value was not inserted

- Specify the key word `DEFAULT` if the column is in the column list

```
INSERT INTO regions (region_id) VALUES (5);
```

```
INSERT INTO regions (region_id, region_name) VALUES (5, DEFAULT);
```

- The result in both cases would be:

REGION_ID	REGION_NAME
-----	-----
5	Anytown

INSERTing a Set of Rows

■ Syntax:

```
INSERT INTO table_name [ (column_list) ]  
SELECT_clause
```

■ Rules:

- The `SELECT` clause can be any legal SQL statement
- The `SELECT` list must map to the column list
- The `SELECT` clause can return 0, 1, or more rows of data
- The other rules are the same as for a single row `INSERT`
- Each `INSERT` statement stands on its own
 - A data integrity violation with 1 row in the set causes the failure of all rows attempting to be inserted

INSERTing a Set of Data: Example

■ Add rows from the location table into countries

- Build the column being used as the primary key using the SUBSTR function

```
INSERT INTO countries (country_id, country_name)
SELECT SUBSTR(location_id, 1, 2), state_province
FROM locations;
```

23 rows created.

INSERT



UPDATE

DELETE

Transactional Control

Chapter Summary

Maintaining Column Values: The `UPDATE` Statement

- The `UPDATE` modifies existing rows in the table
 - Changes the values of column(s) in 0, 1, or more rows in a table
- The columns can be set to static expressions
 - Including literals and functions
- Or to the returned values from a `SELECT` statement

UPDATE Syntax: Static Expressions

- Syntax:

```
UPDATE table_name
SET   column_name = expression | DEFAULT
      [ , ... ]
[ WHERE condition ]
```

- Rules:

- The key word *SET* is mandatory and appears only once
- More than one column can be updated
 - The expression must resolve to the same datatype
- The **DEFAULT** for the column can be referenced
- The **WHERE** clause is optional
 - If not specified, **ALL** rows of the table will be updated

UPDATE Statement: Static Examples

```
UPDATE regions
SET region_name = 'NewTown'
WHERE region_id = 0;
```

```
SELECT *
FROM regions
WHERE region_id = 0;
```

REGION_ID	REGION_NAME
0	NewTown

```
UPDATE regions
SET region_id = 10 - 9
    , region_name = 'SubTown'
WHERE region_id = 0;
```

REGION_ID	REGION_NAME
1	SubTown

DEFAULTs and NULLs

- The `DEFAULT` keyword can be used if one exists
- Be careful with `NULL`
 - `NULL` is a value
 - `'NULL'` is a literal string

```
UPDATE regions
SET   region_name = NULL
WHERE region_id = 0;
```

REGION_ID	REGION_NAME
0	

```
UPDATE regions
SET   region_name = 'NULL'
WHERE region_id = 0;
```

REGION_ID	REGION_NAME
0	NULL

```
UPDATE regions
SET   region_name = DEFAULT
WHERE region_id = 0;
```

REGION_ID	REGION_NAME
0	Anytown

UPDATEing Using Subqueries

- A column can also be set equal to a subquery
- Syntax:

```
UPDATE table_name
SET { (column [, column ]...) = (subquery)
     | [ column = { expr | (subquery) | DEFAULT }
     }
    [ , ... ]
WHERE conditional_clause
```

- Example:

```
UPDATE regions
SET region_name = (
    SELECT country_name
    FROM countries
    WHERE country_id = 'AR'
)
WHERE region_id = 0;
```

```
REGION_ID REGION_NAME
-----
0 Argentina
```

UPDATE Syntax: Rules for Using Subqueries

- The subquery must be placed within parenthesis
- Specify a subquery that returns exactly one row for each row updated
 - If you specify only one column in the *update_set_clause*
 - The subquery can return only one value
- If you specify multiple columns in the *update_set_clause*:
 - The subquery must return as many values as you have specified columns
 - The multiple columns must be placed within parentheses
- If the subquery returns no rows, then the column is assigned a null

```
UPDATE regions
SET (region_id, region_name) = (
    SELECT 10, country_name
    FROM countries
    WHERE country_id = 'AR'
)
WHERE region_id = 0;
```

INSERT

UPDATE



DELETE

Transactional Control

Chapter Summary

Removing Row(s): The DELETE Statement

- Syntax:

```
DELETE FROM table  
[WHERE conditional_clause]
```

- The conditional clause specifies which rows are to be removed
- The WHERE clause is optional
 - If not specified, ALL rows of the table will be deleted

DELETE Examples

- DELETE is a set statement

```
DELETE FROM emp;
```

15 rows deleted.

```
DELETE FROM emp WHERE sal > 6000;
```

1 row deleted.

```
DELETE FROM emp WHERE sal > 1500;
```

8 rows deleted.

```
DELETE FROM emp WHERE sal > 10000;
```

0 rows deleted.

DML and Integrity Constraints

- DELETE, like any DML statement, must conform to integrity constraints
 - Parent rows cannot be removed unless the CASCADE option is set

```
DELETE FROM regions WHERE region_id = 12;
```

```
1 row deleted.
```

```
DELETE FROM regions;
```

```
ERROR at line 1:
```

```
ORA-02292: integrity constraint (HR.COUNTR_REG_FK)  
violated - child record found
```

INSERT

UPDATE

DELETE



Transactional Control

Chapter Summary

Transactional Control

- A transaction is a logical unit of work that comprises one or more SQL statements
 - Must leave the database in a consistent state
- Data Manipulation Language (DML) statements are part of transactional control
 - The ability to make one or more changes to data as a group
 - The first SQL statement automatically starts a transaction
- To make the changes permanent to the database, issue a `COMMIT`
- To undo the changes, issue a `ROLLBACK`
 - `ROLLBACK` will undo all changes to the beginning of the transaction

Undoing Part of a Transaction

- `SAVEPOINTS` can be issued as part of the transactional stream
 - A savepoint is an intermediate marker that divides the transaction into smaller pieces
- Syntax:

```
SAVEPOINT my_savepoint;
```
- Allows the session to roll back any DML that was issued after the establishment of the savepoint

ROLLBACK Example

- First, set up the transactional activity in the session

```
INSERT INTO jobs VALUES ('IT_AC_PROG', 'Accounting Programmer', 1000, 1000);  
1 row created.
```

```
SAVEPOINT sp_jobs;  
Savepoint created.
```

```
INSERT INTO jobs VALUES ('IT_PR_PROG', 'Payroll Programmer', 999999, 999999);  
1 row created.
```

```
SELECT job_id, job_title FROM jobs WHERE job_id IN ('IT_AC_PROG', 'IT_PR_PROG');
```

```
JOB_ID      JOB_TITLE
```

```
-----
```

```
IT_AC_PROG Accounting Programmer
```

```
IT_PR_PROG Payroll Programmer
```

ROLLBACK Example (continued)

- Undo the second insert by issuing the ROLLBACK TO... statement

```
ROLLBACK TO sp_jobs;  
Rollback complete.
```

- Test for the existence of the inserts

JOB_ID	JOB_TITLE
IT_AC_PROG	Accounting Programmer

- Note that the Payroll programmer no longer exists
- Finally, issue a session-level rollback and test again

```
ROLLBACK;  
Rollback complete.
```

```
SELECT job_id, job_title FROM jobs WHERE job_id IN ('IT_AC_PROG', 'IT_PR_PROG');  
no rows selected
```

- Where did we roll back to?

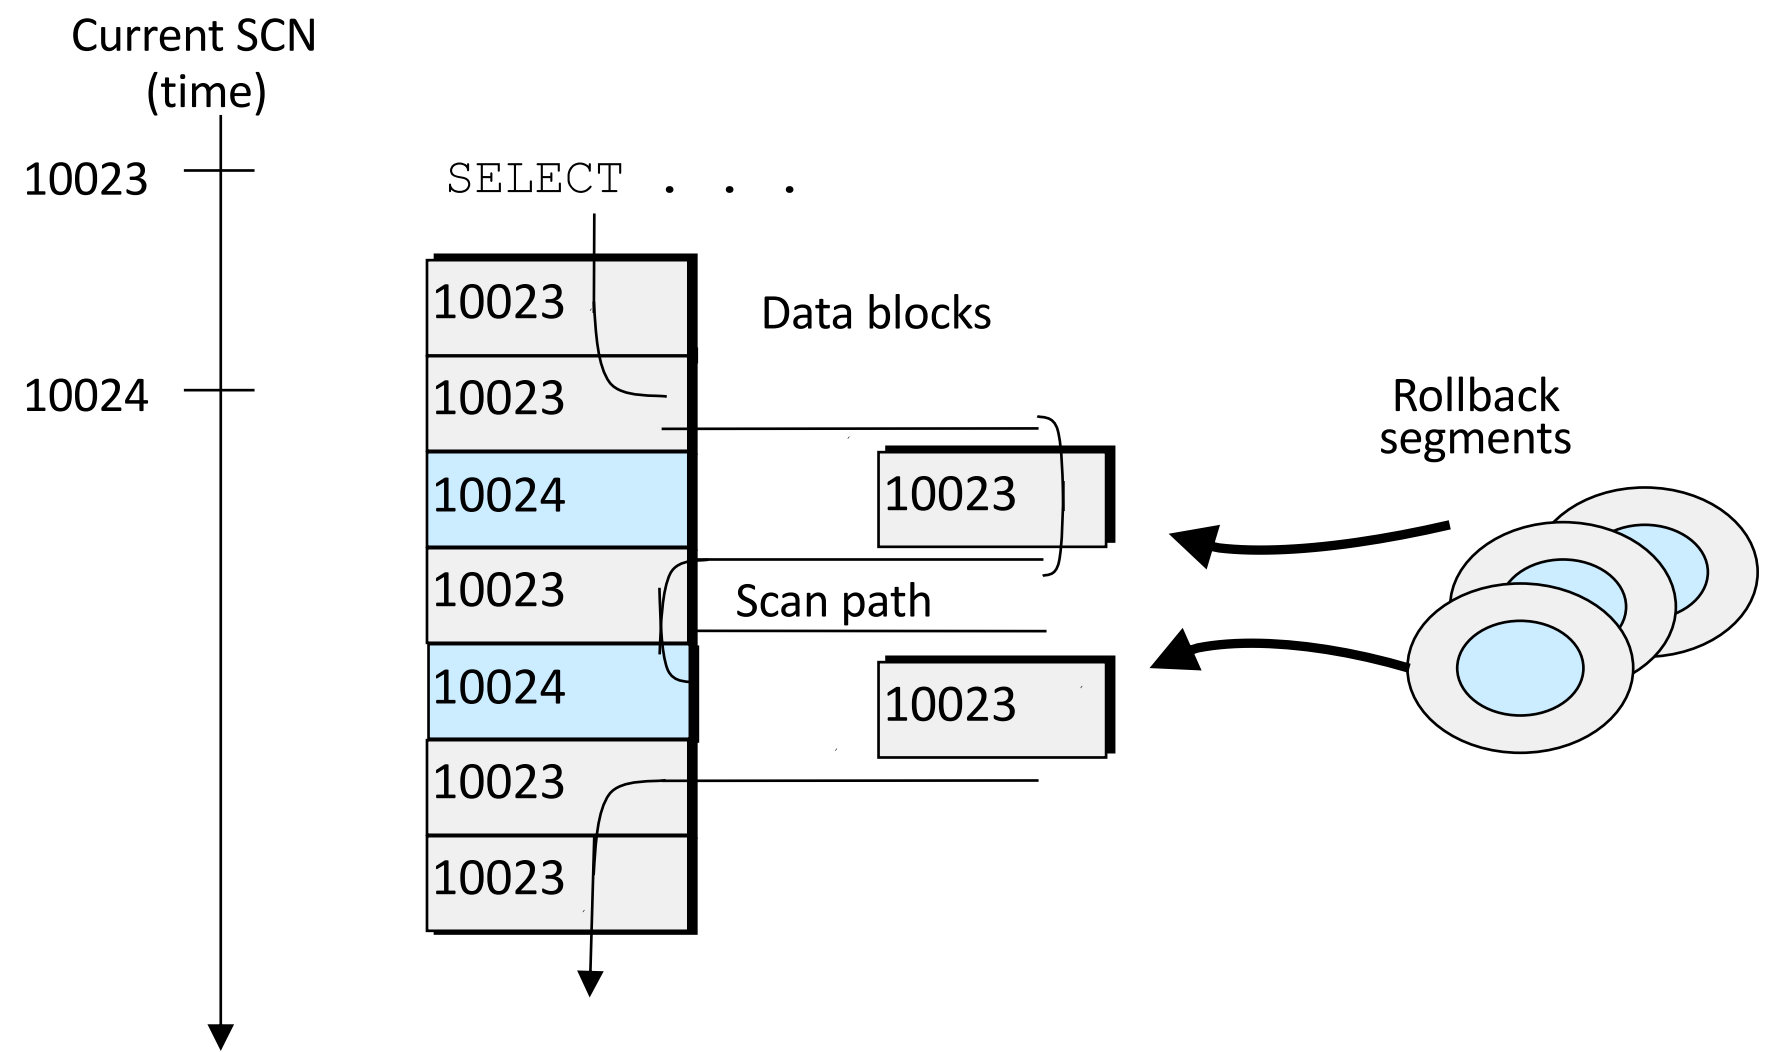
Read Consistency

- Statement-level read consistency
 - A single select is not impacted by database changes during its execution
- Transaction-level read consistency
 - Multiple selects in one transaction are not impacted by database changes
 - Enforced by `SET TRANSACTION READ ONLY`
- Ensures that readers of database data do not have to wait for writers or readers of the same data
- Ensures that writers of database data do not have to wait for readers of the same data
- Automatically enforced by Oracle using rollback segments and the System Change Number (SCN)

Read Consistency Example

- As query enters execution stage, the current SCN is determined
 - SCN is a sequential number assigned to each transaction
 - Recorded in rollback segment and redo log
 - In this example, SCN is 10023
- As query reads, only data blocks with SCN 10023 are used
- Blocks with changed data (more recent SCNs) are reconstructed using the rollback segments
- Reconstructed data is returned to the query
- Query returns all committed data with respect to the SCN recorded at the time execution started
- Committed or uncommitted changes of other transactions that occur during a query's execution are not seen, guaranteeing that a consistent set of data is returned for each query

Read Consistency Example (continued)



SET TRANSACTION READ ONLY Command

- First statement of a read-only transaction
- Remaining statements in the transaction cannot change data in the database
 - Insert, update, and delete are not allowed
- Last statement in the read-only transaction is `COMMIT` or `ROLLBACK`

Example of Transaction-Level Read Consistency

- Provide two listings of the `employees` table ordered by:
 - Last name
 - Hire date
- Transaction begins
 - Establish read consistency for the transaction
`SET TRANSACTION READ ONLY;`
 - Create the employee listing by last name
`SELECT *
FROM employees
ORDER BY last_name;`
 - Create the employee listing by hire date
`SELECT *
FROM employees
ORDER BY hire_date;`
 - Release read consistency
`ROLLBACK;`

Advantages and Disadvantages of Read Consistency

- Advantage:
 - Allows consistency of data throughout a transaction
- Disadvantages:
 - Performance of `SELECT` statements is impacted due to the necessity of reconstructing data from `ROLLBACK` segments
 - May get `snapshot too old` error

transaction isolation levels

- **dirty reads** - when transaction reads data that has not been committed

- **norepeatable reads**- read occur when a transaction reads a row twice but gets different data each time

-**phantoms** - a row that matches search criteria but is not initially seen

read committed
read uncommitted
repeatable
serializable

Serializable Transactions

- Makes it appear as if there are no other users modifying data in the database
 - Any row read in the transaction is assured to be the same upon a reread
 - Most restrictive form of transaction isolation
- Set at the transaction level with the command
`SET TRANSACTION ISOLATION LEVEL SERIALIZABLE`
 - Causes all statements in the transaction, including `SELECT`, to obtain locks
- Improvement over `SET TRANSACTION READ ONLY` because it allows `UPDATE`, `INSERT`, and `DELETE`
- *Beware*—high cost to performance

Exercise 7.1: Manipulating Data



60 min

- Please complete this exercise in your Exercise Manual

INSERT

UPDATE

DELETE

Transactional Control



Chapter Summary

Chapter Summary

In this chapter, we have discussed:

- Manipulating data using DML:
 - INSERT
 - UPDATE
 - DELETE
- Transactional Control Statements:
 - COMMIT
 - ROLLBACK
 - SAVEPOINT